ANTHROPIC, BIG TECH, AND THE PENTAGON

Corporate Structure, Financial Entanglement, and the Department of Defense Conflict

2021–2026

Research compiled March 12, 2026

All claims sourced and linked. See Sources section for full bibliography.

ABSTRACT: It covers everything chronologically across six parts: Anthropic's founding and PBC/LTBT governance structure (2021), Google and Amazon's investment history including stake sizes, caps, and the absence of board seats or voting rights, the Microsoft/Nvidia entry in late 2025, the Palantir partnership and classified network deployment (November 2024), the $200M DOD contract (July 2025), the full 2026 conflict from the Hegseth memorandum through the autonomous weapons red lines, the Venezuela incident, the February 27 presidential ban, the supply chain risk designation, OpenAI's same-day classified network deal, Anthropic's March 9 lawsuits, Microsoft's intervention, the $80 billion cloud commitment breakdown, and a closing analysis of the five core structural tensions. Every factual claim is sourced with a working link in the bibliography at the end.

Part I: Founding and Corporate Structure (2021)

The OpenAI Exodus

Between December 2020 and January 2021, a group of senior OpenAI researchers and executives departed the organization over fundamental disagreements about the balance between AI capability development and safety research. The departures were led by Dario Amodei, then OpenAI's Vice President of Research, and his sister Daniela Amodei, then OpenAI's Vice President of Safety and Policy. They were joined by approximately nine other employees, including Tom Brown (a lead author of GPT-3), Jared Kaplan, Sam McCandlish, Jack Clark, Chris Olah, and Benjamin Mann.

Anthropic was formally registered in California on February 3, 2021. Dario Amodei later told Fortune that he left OpenAI because the organization was focused on making ever-more-powerful models without a commensurate focus on safety. The new company raised $124 million in initial funding and established its headquarters in San Francisco.

Public Benefit Corporation Structure

Anthropic was incorporated as a Public Benefit Corporation (PBC) under Delaware law. This structure imposes a dual fiduciary obligation on the board: directors must work to increase shareholder profits, but they are also legally empowered to prioritize the company's stated mission of ensuring that transformative AI helps people and society flourish. This was a deliberate structural choice intended to prevent the safety mission from being overridden by pure profit maximization—a concern the founders felt had materialized at OpenAI.

The Long-Term Benefit Trust (LTBT)

In September 2023, Anthropic announced the creation of its Long-Term Benefit Trust, a purpose trust whose sole mandate is the responsible development and maintenance of advanced AI for the long-term benefit of humanity. The Trust holds a special class of stock—Class T shares—that grants it the authority to elect an increasing number of directors to Anthropic's board over time, phasing in according to both time-based and funding-based milestones.

The mechanism works as follows: the LTBT initially elected one of five board members. By mid-2025, that number was set to rise to two of five, and by late 2025, to three of five—giving the Trust majority control of the board within approximately four years of its creation. The Trust's first five members were selected for expertise in AI safety, national security, and social enterprise. They included Jason Matheny (CEO of the RAND Corporation), Kanika Bahl (CEO of Evidence Action), Paul Christiano (AI safety researcher), Zach Robinson, and Neil Buddy Shah (later replaced by Richard Fontaine as of October 2025). Trustees serve one-year terms, and future trustees are elected by the sitting trustees themselves.

Critically, a supermajority of stockholders retains the power to amend the Trust's governing documents—a failsafe intended to protect investor interests against trustee overreach. The

structure was designed so that neither Google, Amazon, nor any single investor could gain effective control of the company, even after an IPO.

Board of Directors

As of early 2026, Anthropic's board includes Dario Amodei (CEO), Daniela Amodei (President), Yasmin Razavi, Jay Kreps, Reed Hastings, and Chris Liddell. Notably, neither Amazon nor Google holds a board seat or board observer rights, despite being the company's two largest financial backers.

Part II: Google and Amazon—Investment, Infrastructure, and Influence

Google's Investment History

Google's relationship with Anthropic began with an early investment and deepened steadily:

- **2022–2023:** Google invested approximately $2 billion in Anthropic during 2023, making it one of the company's earliest major strategic backers.
- **2024–2025:** Google invested an additional $1 billion, bringing its total committed capital to over $3 billion. Legal filings submitted as part of a Google antitrust case revealed that Google holds a 14% stake in Anthropic.
- **Cap and restrictions:** Google's ownership is contractually capped at 15%. Google holds no voting rights, no board seats, and no board observer rights. The DOJ initially proposed that Google divest its AI stakes (including in Anthropic) as part of antitrust remedies, but later dropped that proposal.
- **October 2025 cloud deal:** Google and Anthropic announced a cloud infrastructure deal worth tens of billions of dollars, giving Anthropic access to up to one million of Google's custom-designed Tensor Processing Units (TPUs), expected to deliver over a gigawatt of compute capacity by 2026.
- **Planned additional investment:** Google aimed to invest an additional $750 million in Anthropic through a convertible debt agreement in September 2025.

At a $380 billion valuation (February 2026 estimate), Google's 14% stake would be worth approximately $53 billion—a substantial return on roughly $3 billion invested.

Amazon's Investment History

Amazon's investment in Anthropic is the largest by any single entity:

- **September 2023:** Amazon made an initial $1.25 billion investment, marking its entry as a major Anthropic backer.
- **March 2024:** Amazon invested an additional $2.75 billion—at the time, the largest outside investment in Amazon's three-decade history—bringing its total to $4 billion.
- **November 2024:** Amazon committed another $4 billion, bringing total investment to $8 billion. Amazon retained its position as a minority investor.
- **2025 and beyond:** Reports indicated Amazon was considering additional multibillion-dollar investments. Amazon's stake remained capped below 33% with no voting rights and no board seat.

As part of the investment relationship, Amazon Web Services became Anthropic's primary cloud and training partner. Amazon developed Project Rainier, a high-performance AI cluster using hundreds of thousands of custom Trainium2 UltraServer chips, built exclusively for Anthropic's use. In Q3 2024, Amazon reported a $9.5 billion pretax gain tied to the rising valuation of its Anthropic investment.

Microsoft and Nvidia Enter (November 2025)

In November 2025, Anthropic announced parallel strategic partnerships with Microsoft and Nvidia. Microsoft committed to investing up to $5 billion in Anthropic, while Nvidia committed up to $10 billion. As part of the deal, Anthropic committed to purchasing $30 billion in Azure compute capacity from Microsoft and contracted for additional capacity up to one gigawatt with Nvidia's Grace Blackwell and Vera Rubin systems. These partnerships pushed Anthropic's valuation to approximately $350 billion and made Claude the only frontier AI model available across all three major cloud platforms (AWS, Azure, and Google Cloud).

Governance Independence Despite Financial Dependence

Despite over $16 billion in direct investment from Amazon, Google, Microsoft, and Nvidia, Anthropic's governance structure kept all strategic investors at arm's length. Neither Amazon (the largest investor) nor Google holds board seats, voting rights, or observer rights. The LTBT was designed to accumulate board control over time specifically to prevent investor capture of the safety mission. However, this structural independence exists alongside deep financial and operational dependence on these same investors' cloud infrastructure—a tension that would become critically important in 2026.

Part III: The Department of Defense Relationship

First Contact and the Palantir Partnership (November 2024)

On November 7, 2024, Anthropic, Palantir Technologies, and Amazon Web Services jointly announced a partnership to provide U.S. intelligence and defense agencies with access to Claude 3 and Claude 3.5 models. The arrangement ran Claude through Palantir's AI Platform (AIP) on AWS's government-accredited cloud infrastructure.

Palantir became the first industry partner to deploy Claude on classified networks, specifically at Impact Level 6 (IL6)—a Department of Defense accreditation level reserved for systems handling data classified up to one level below Top Secret. This gave intelligence and defense analysts access to Claude for processing and analyzing classified data, intelligence reports, and operational planning.

The $200 Million DOD Contract (July 2025)

In July 2025, Anthropic secured a defense contract worth up to $200 million, significantly expanding its footprint within the Department of Defense. The company stated that it had accelerated mission impact across U.S. defense workflows with partners like Palantir, where Claude was integrated into mission workflows on classified networks, enabling defense and intelligence organizations to rapidly process and analyze vast amounts of complex data.

By this point, Anthropic was the only AI model company officially deployed across the Pentagon's classified systems—a position of singular importance and vulnerability.

Part IV: The 2026 Conflict—Autonomous Weapons, Surveillance, and the Federal Ban

The Hegseth Memorandum (January 2026)

In January 2026, Defense Secretary Pete Hegseth issued an AI strategy memorandum directing all Department of Defense AI contracts to incorporate standard “any lawful use” language within 180 days. This directly contradicted the terms of Anthropic’s existing contract, which contained restrictions on how the military could deploy Claude.

Anthropic’s position was built on two red lines. First, that current frontier AI models are not reliable enough for use in fully autonomous weapons, and allowing such use would endanger American warfighters and civilians. Second, that mass domestic surveillance of Americans constitutes a violation of fundamental rights. Anthropic CEO Dario Amodei stated publicly that crossing those lines was contrary to American values.

The Venezuela/Maduro Incident

Tensions intensified in connection with the U.S. military’s operation to capture Venezuelan leader Nicolás Maduro. According to reporting by NBC News and Semafor, the Pentagon used Claude (through the Palantir integration) in operational planning related to the raid. An Anthropic executive reportedly reached out to a Palantir executive to inquire whether Claude had been used in the operation, in a manner that implied disapproval. Anthropic denied this characterization, saying the company had not discussed this with or expressed concerns to any industry partners outside of routine technical discussions. The incident nevertheless deepened the rift between Anthropic and the Pentagon.

The Deadline and Presidential Ban (February 27, 2026)

On February 24, 2026, Defense Secretary Hegseth gave Anthropic a deadline: agree to allow unrestricted use of Claude for all legal purposes by 5:01 p.m. on Friday, February 27. Anthropic released a public statement indicating it would not change its position.

On February 27, President Trump directed all federal agencies to immediately cease using Anthropic’s technology. Simultaneously, Secretary Hegseth designated Anthropic as a supply chain risk to national security. This designation—traditionally used against foreign adversaries—was applied to an American company for what is believed to be the first time. Chief Pentagon Spokesman Sean Parnell stated that the nation requires partners willing to help warfighters win in any fight.

The Supply Chain Risk Designation (March 5–6, 2026)

On March 5, the DOD officially notified Anthropic’s leadership that the company and its products had been formally designated a supply chain risk, effective immediately. The designation required all defense vendors and contractors to certify that they do not use Anthropic’s models in their work with the Pentagon.

An internal Pentagon memo dated March 6 ordered military commanders to remove Anthropic's AI products from their systems within 180 days. Reporting by CBS News revealed the scope of the removal would affect national security systems including those for nuclear weapons, ballistic missile defense, and cyber warfare.

OpenAI's Entry Into Classified Networks (February 28, 2026)

Hours after the presidential ban, OpenAI announced it had reached a deal with the Department of Defense to provide its own AI technology for classified networks. CEO Sam Altman acknowledged the negotiations had been "definitely rushed." OpenAI's agreement included stated prohibitions on domestic surveillance of American citizens, directing autonomous weapons, and high-stakes automated decision-making—language that MIT Technology Review noted was substantively similar to Anthropic's own positions, but framed as voluntary commitments rather than contractual restrictions.

The timing drew sharp criticism. CNN reported that some OpenAI employees expressed respect for Anthropic's stance and frustration with OpenAI's handling of the situation. Reports indicated ChatGPT uninstall rates surged approximately 200% in the days following the announcement, as some users shifted to Claude in protest.

Anthropic's Federal Lawsuits (March 9, 2026)

On March 9, 2026, Anthropic filed two federal lawsuits against the Trump administration:

- **California federal court:** Challenging the legality of the president's directive for federal agencies to cease using Anthropic's technology and the supply chain risk designation.
- **D.C. Circuit Court of Appeals:** Challenging the procedural and constitutional basis of the Pentagon's actions.

Anthropic's legal arguments rested on several grounds. The company argued that it was denied adequate due process, that the government was retaliating against the company for exercising First Amendment-protected speech (its public positions on autonomous weapons and surveillance), and that the president lacked the authority to direct agencies to cease using a private company's technology on these grounds. CEO Dario Amodei wrote in a blog post that Anthropic did not believe the action was legally sound and saw no choice but to challenge it in court. The company sought injunctive relief, alleging that hundreds of millions of dollars in current and future contracts were in jeopardy.

Microsoft's Intervention (March 10, 2026)

On March 10, Microsoft filed a brief supporting Anthropic, asking the court to issue a temporary restraining order blocking the Pentagon's supply chain risk designation for all existing contracts. Microsoft argued this would enable a more orderly transition and avoid disrupting the American military's ongoing use of advanced AI. Amazon separately confirmed that Claude remained available for all commercial AWS customers outside of defense work.

Part V: The Financial Entanglement

$80 Billion in Cloud Commitments Through 2029

According to Data Center Dynamics, Anthropic's most optimistic internal forecasts project that the company will pay Amazon, Google, and Microsoft at least $80 billion in cloud computing costs through 2029 to run Claude on their servers. This figure breaks down across the three providers:

- **Amazon Web Services:** As Anthropic's primary cloud partner and the host of its classified government deployments, AWS captures the largest share. Amazon also built the dedicated Project Rainier cluster for Anthropic using custom Trainium2 chips.
- **Microsoft Azure:** Anthropic committed to purchasing $30 billion in Azure compute capacity as part of the November 2025 deal.
- **Google Cloud:** The October 2025 deal for up to one million TPUs was valued at tens of billions of dollars, with industry estimates for 1 GW of data center capacity running approximately $50 billion.

Beyond direct cloud costs, Anthropic also faces billions in revenue-sharing obligations with its cloud distribution partners. The Information reported that Anthropic expects to pay up to $6.4 billion in resale partner payouts in 2027 alone—roughly one-tenth of its projected revenue.

Revenue and Valuation Trajectory

Anthropic's financial growth has been extraordinary:

- **March 2025:** $3.5 billion Series E at $61.5 billion valuation, led by Lightspeed Venture Partners with participation from Bessemer, Cisco, D1 Capital, Fidelity, General Catalyst, Jane Street, Menlo Ventures, and Salesforce Ventures.
- **July 2025:** Reported in talks to raise capital at $170 billion valuation.
- **September 2025:** Valuation reached $183 billion following a new round led by ICONIQ.
- **November 2025:** Microsoft and Nvidia partnerships pushed valuation to approximately $350 billion.
- **February 2026:** Estimated valuation of $380 billion. Annual revenue run rate approaching $7 billion, with Claude powering over 300,000 businesses.

The Structural Paradox

Anthropic's financial structure creates a fundamental paradox. The company's governance was designed to insulate its safety mission from investor pressure—neither Amazon nor Google holds voting rights or board seats. But the company's operations are almost entirely dependent on infrastructure owned by its strategic investors. Approximately $80 billion in committed cloud payments flow to the same companies that own the largest equity stakes. If Anthropic's revenue were disrupted—as the federal ban threatens to do—its ability to meet these cloud

commitments could be jeopardized, creating potential leverage for infrastructure providers that the governance structure was not designed to address.

Part VI: Summary of Structural Tensions

1. Governance Independence vs. Operational Dependence

Anthropic's LTBT and PBC structure successfully prevent any single investor from controlling the company's board or overriding its safety mission through formal governance channels. But the company's cloud infrastructure is wholly owned by its investors. This means that while Amazon and Google cannot dictate Anthropic's policies through board votes, they hold structural leverage through the infrastructure on which Anthropic's entire business depends.

2. Safety Mission vs. Defense Revenue

Anthropic's entry into defense and intelligence work was motivated by a belief that responsible AI companies should shape how AI is used in national security, rather than ceding the field to less safety-conscious competitors. But the deeper Anthropic embedded itself in classified networks, the more leverage the Pentagon gained—and the more costly it became to say no. The $200 million contract and integration into systems for nuclear weapons and missile defense made Anthropic's refusal to remove safety guardrails not just a policy disagreement but a national security confrontation.

3. Precedent and the Supply Chain Risk Weapon

The designation of an American company as a supply chain risk—a tool previously reserved for foreign adversaries—established a new precedent for government retaliation against companies whose policies conflict with executive branch preferences. The Electronic Frontier Foundation noted that the case raises fundamental questions about whether privacy protections should depend on the decisions of a few powerful people.

4. The Competitor Dynamic

OpenAI's immediate entry into classified networks after Anthropic's ban demonstrated the competitive dynamics at play. The Pentagon's message was clear: companies that cooperate will be rewarded; those that resist will be replaced. OpenAI's deal included many of the same prohibitions Anthropic sought, but framed as voluntary commitments—a distinction that may matter more in terms of enforcement leverage than actual operational restrictions.

5. The $80 Billion Question

If the federal ban significantly damages Anthropic's revenue—through lost government contracts, chilled commercial demand, or contractor avoidance of a blacklisted company—the company's ability to meet $80 billion in cloud commitments could come under pressure. This would directly affect Amazon, Google, and Microsoft's expected cloud revenue, creating a scenario where the companies that invested in Anthropic's independence are financially affected by the consequences of that independence being exercised.

Sources

Founding and Corporate Structure

Anthropic Company Page: https://www.anthropic.com/company

Anthropic Wikipedia: https://en.wikipedia.org/wiki/Anthropic

Contrary Research — Anthropic Business Breakdown: https://research.contrary.com/company/anthropic

Anthropic Long-Term Benefit Trust Announcement: https://www.anthropic.com/news/the-long-term-benefit-trust

Harvard Law — Anthropic LTBT Analysis: https://corpgov.law.harvard.edu/2023/10/28/anthropic-long-term-benefit-trust/

TIME — How Anthropic Designed Itself: https://time.com/6983420/anthropic-structure-openai-incentives/

Who Owns Anthropic (2026): https://www.revenuememo.com/p/who-owns-anthropic

RankRed — Anthropic Ownership Details: https://www.rankred.com/who-owns-anthropic/

Google and Amazon Investments

Google 14% Stake Revealed — Computing.co.uk: https://www.computing.co.uk/news/2025/ai/google-revealed-to-hold-14-stake-in-anthropic

Google Funding Filings — TechCrunch: https://techcrunch.com/2025/03/11/google-has-given-anthropic-more-funding-than-previously-known-show-new-filings/

Amazon $4 Billion Investment — CNBC: https://www.cnbc.com/2024/11/22/amazon-to-invest-another-4-billion-in-anthropic-openais-biggest-rival.html

Amazon $8 Billion Total — TechFundingNews: https://techfundingnews.com/amazon-anthropic-ai-investment-strategy/

Spyglass — Amazon's Anthropic Stake: https://spyglass.org/amazon-anthropic-stake-google/

Alphabet and Amazon Gains — TradeAlgo: https://news.tradealgo.com/news/alphabet-and-amazon-see-billions-in-gains-from-anthropic-investments

Microsoft, Nvidia, and Cloud Deals

Microsoft/Nvidia/Anthropic Partnerships — Microsoft Blog: https://blogs.microsoft.com/blog/2025/11/18/microsoft-nvidia-and-anthropic-announce-strategic-partnerships/

Anthropic $30B Azure — DCD: https://www.datacenterdynamics.com/en/news/anthropic-to-purchase-30bn-in-microsoft-azure-credits-nvidia-and-microsoft-to-invest-in-ai-company/

Microsoft Invest $5B — GeekWire: https://www.geekwire.com/2025/microsoft-to-invest-5b-in-anthropic-as-claude-maker-commits-30b-to-azure-in-new-nvidia-alliance/

Google Cloud TPU Deal — CNBC: https://www.cnbc.com/2025/10/23/anthropic-google-cloud-deal-tpu.html

1 Million TPUs — Google Cloud Press: https://www.googlecloudpresscorner.com/2025-10-23-Anthropic-to-Expand-Use-of-Google-Cloud-TPUs-and-Services

$80B Cloud Spend — DCD: https://www.datacenterdynamics.com/en/news/anthropic-cloud-spend-expected-to-reach-80bn-through-2029/

$6.4B Revenue Share — The Information: https://www.theinformation.com/articles/anthropic-sweetens-deal-cloud-providers

Department of Defense and Palantir

Palantir Partnership Announcement: https://investors.palantir.com/news-details/2024/Anthropic-and-Palantir-Partner-to-Bring-Claude-AI-Models-to-AWS-for-U.S.-Government-Intelligence-and-Defense-Operations/

Anthropic DOD Partnership — Anthropic: https://www.anthropic.com/news/anthropic-and-the-department-of-defense-to-advance-responsible-ai-in-defense-operations

Anthropic-Palantir-AWS — TechCrunch: https://techcrunch.com/2024/11/07/anthropic-teams-up-with-palantir-and-aws-to-sell-its-ai-to-defense-customers/

The 2026 Conflict

Semafor — Palantir at Heart of Rift: https://www.semafor.com/article/02/17/2026/palantir-partnership-is-at-heart-of-anthropic-pentagon-rift

NBC News — Tensions Boiling Point: https://www.nbcnews.com/tech/security/anthropic-ai-defense-war-venezuela-maduro-rcna259603

Axios — Pentagon Threatens Cutoff: https://www.axios.com/2026/02/15/claude-pentagon-anthropic-contract-maduro

The Hill — Pentagon Reviewing Partnership: https://thehill.com/policy/defense/5740369-pentagon-anthropic-relationship-review/

ASIS — Anthropic Refuses Guardrail Removal: https://www.asisonline.org/security-management-magazine/latest-news/today-in-security/2026/february/Anthropic-Refusal/

CNBC — Supply Chain Risk and Iran: https://www.cnbc.com/2026/03/05/anthropic-pentagon-ai-claude-iran.html

CBS News — Pentagon Removal Memo: https://www.cbsnews.com/news/pentagon-ai-anthropic-memo-remove-from-key-systems/

Federal Ban, Lawsuits, and Aftermath

NPR — OpenAI Pentagon Deal After Ban: https://www.npr.org/2026/02/27/nx-s1-5729118/trump-anthropic-pentagon-openai-ai-weapons-ban

CNBC — OpenAI Strikes Pentagon Deal: https://www.cnbc.com/2026/02/27/openai-strikes-deal-with-pentagon-hours-after-rival-anthropic-was-blacklisted-by-trump.html

CNN — OpenAI Staff Fuming: https://www.cnn.com/2026/03/04/tech/pentagon-anthropic-openai-staff-reactions

MIT Technology Review — OpenAI Compromise: https://www.technologyreview.com/2026/03/02/1133850/openais-compromise-with-the-pentagon-is-what-anthropic-feared/

NPR — Anthropic Sues Trump Administration: https://www.npr.org/2026/03/09/nx-s1-5742548/anthropic-pentagon-lawsuit-amodai-hegseth

CNBC — Anthropic Banned, Experts Worried: https://www.cnbc.com/2026/03/09/anthropic-was-the-pentagons-choice-for-ai-now-its-banned-and-experts-are-worried.html

CNN — Anthropic Sues Pentagon: https://www.cnn.com/2026/03/09/tech/anthropic-sues-pentagon

Euronews — Anthropic Sues Over Label: https://www.euronews.com/next/2026/03/09/anthropic-sues-us-defense-department-over-supply-chain-risk-label

CNBC — Microsoft Seeks Restraining Order: https://www.cnbc.com/2026/03/10/microsoft-says-court-should-temporarily-block-pentagon-ban-anthropic.html

EFF — Privacy Protections Analysis: https://www.eff.org/deeplinks/2026/03/anthropic-dod-conflict-privacy-protections-shouldnt-depend-decisions-few-powerful

Amazon Confirms Claude for Commercial AWS: https://www.cnbc.com/2026/03/06/amazon-aws-anthropic-claude-pentagon-blacklist.html

TechPolicy.Press — Full Timeline: https://www.techpolicy.press/a-timeline-of-the-anthropic-pentagon-dispute/

Funding Rounds and Valuation

Anthropic Series E — Anthropic: https://www.anthropic.com/news/anthropic-raises-series-e-at-usd61-5b-post-money-valuation

$3.5B Raise — TechCrunch: https://techcrunch.com/2025/03/03/anthropic-raises-3-5b-to-fuel-its-ai-ambitions/

$183B Valuation — Crunchbase: https://news.crunchbase.com/venture/generative-ai-anthropic-funding-iconiq/

$350B Valuation — CNBC: https://www.cnbc.com/2025/11/18/anthropic-ai-azure-microsoft-nvidia.html